AI编译器系列

AI 编译器的前世

ZOMI

Talk Overview

1. 传统编译器
   - History of Compiler - 编译器的发展
   - GCC process and principle – GCC 编译过程和原理
   - LLVM/Clang process and principle – LLVM 架构和原理
2. **AI编译器**
   - Why need AI Compiler – 为什么需要AI编译器
   - Base Common architecture – AI编译器的出现和架构
   - Different and challenge of the future – 未来的挑战与思考

计算机架构的新黄金时代

- A New Golden Age for Computer Architecture: History, Challenges and Opportunities
  https://www.youtube.com/watch?v=kFT54hO1X8M

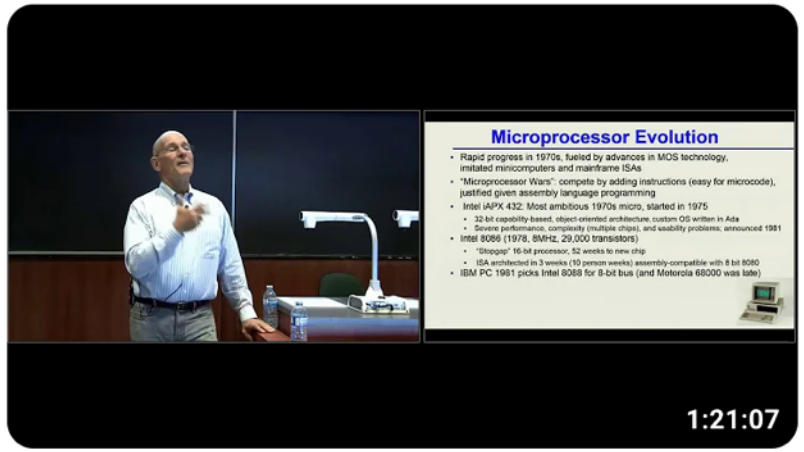

David Patterson - A New Golden Age for Computer Architecture: History, Challenges and Opportunities

7.1万次观看 · 3年前

UBC Computer Science

Abstract: In the 1980s, Mead and Conway democratized chip design and high-level language programming surpassed assembly ...

Turing Awards | What is Computer Architecture | IBM System360 | Semiconductors | Microprocessor... 44 个章节

www.hiascend.com
www.mindspore.cn

编译器的黄金时代

- The Golden Age of Compiler Design in an Era of HW/SW Co-design
- https://www.youtube.com/watch?v=4HgShra-KnY

ASPLOS Keynote: The Golden Age of Compiler Design in an Era of HW/SW Co-design by Dr. Chris Lattner

2.7万次观看 · 1年前

SiFiveInc

This week at the ASPLOS 2021 conference, Dr. Chris Lattner gave the keynote address to open the event with a discussion of the ...

A New Golden Age for Computer Architecture John L. Hennessy, David A. Patterson June 2018 End o... 22 个章节

www.hiascend.com
www.mindspore.cn

为什么需要 AI 编译器

www.hiascend.com
www.mindspore.cn

Hardware is getting harder

- Modern compute acceleration platforms are multi-level and explicit:
  - scalar, Vector, Multi-core, Multi-package, Multi-rack, et. al.
  - None-coherent memory subsystems increase efficiency
- Heterogenous compute incorporating domain-specific accelerators
  - Standard in high-end SoCs, domain-specific hard blocks in FPGAs.
- Many Accelerator IPs are configurable:
  - Optional extensions, tile count, memory hierarchy, etc.

www.hiascend.com
www.mindspore.cn

Hardware is getting harder

- How can developers write Software for this in the first place?
- How can developers afford to build generation-specific Software?

www.hiascend.com
www.mindspore.cn

Next-Gen compilers are needed!

What we need:

- Hardware abstraction spanning diverse accelerators.
- Support for heterogeneous compute platforms.
- Domain specific languages and programming models.
- Quality, reliability and scalability of infrastructure.

www.hiascend.com
www.mindspore.cn

C Compilers leading into the early 90s

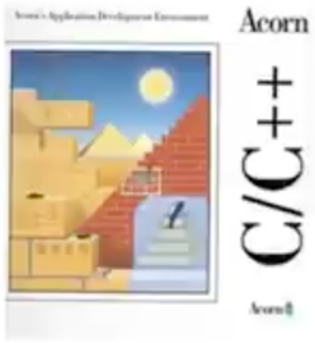

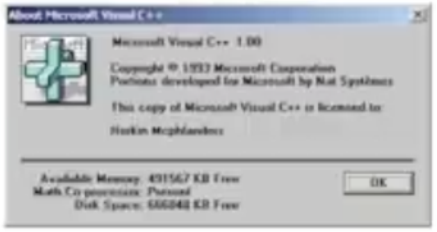

www.hiascend.com
www.mindspore.cn

Three Phase Compiler Stage

www.hiascend.com
www.mindspore.cn

LLVM Enables Collaboration & Reuse

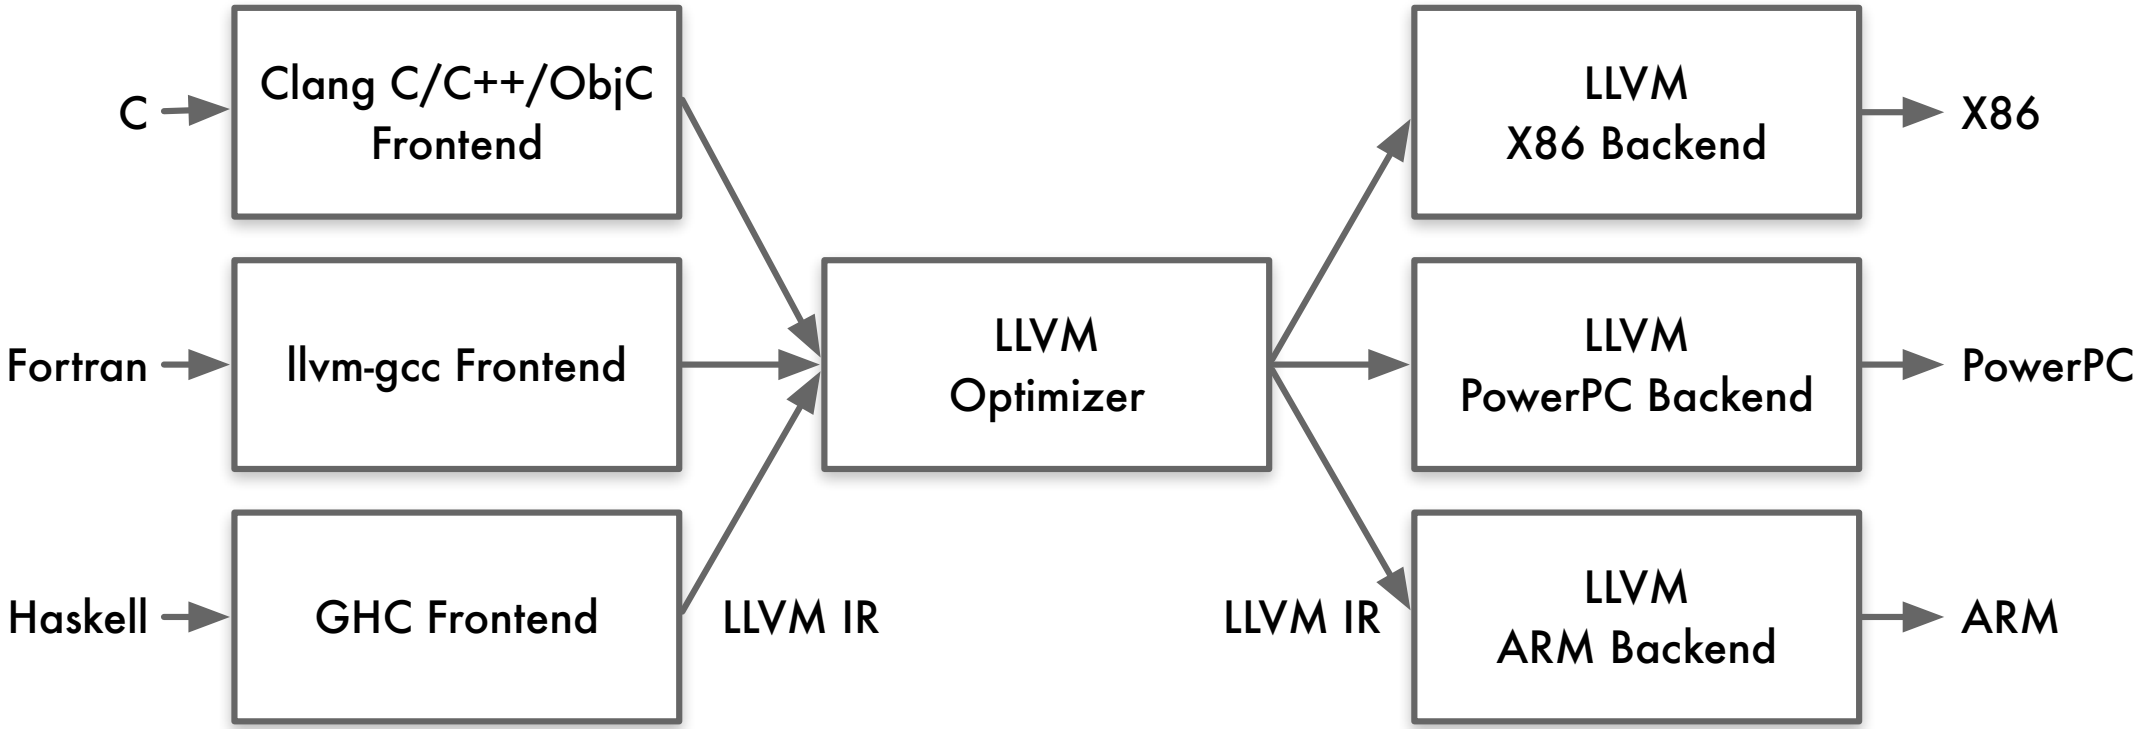

www.hiascend.com
www.mindspore.cn

It's happening!

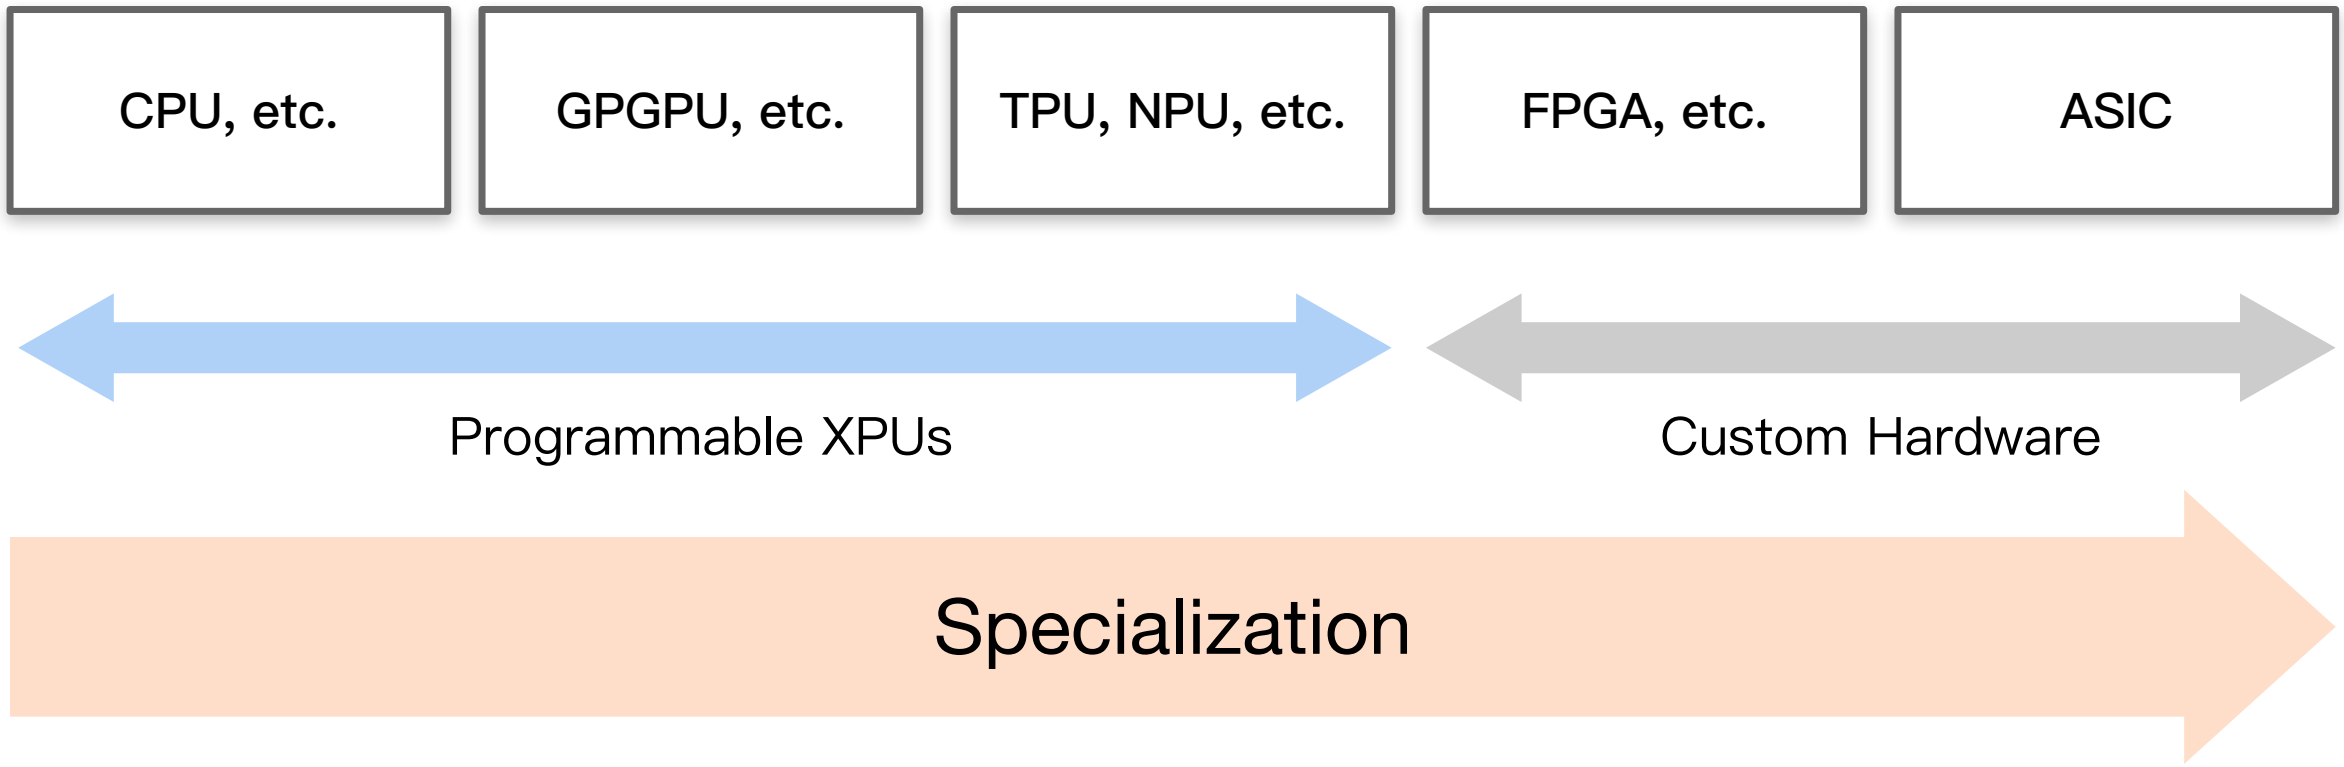

www.hiascend.com
www.mindspore.cn

Los of players!

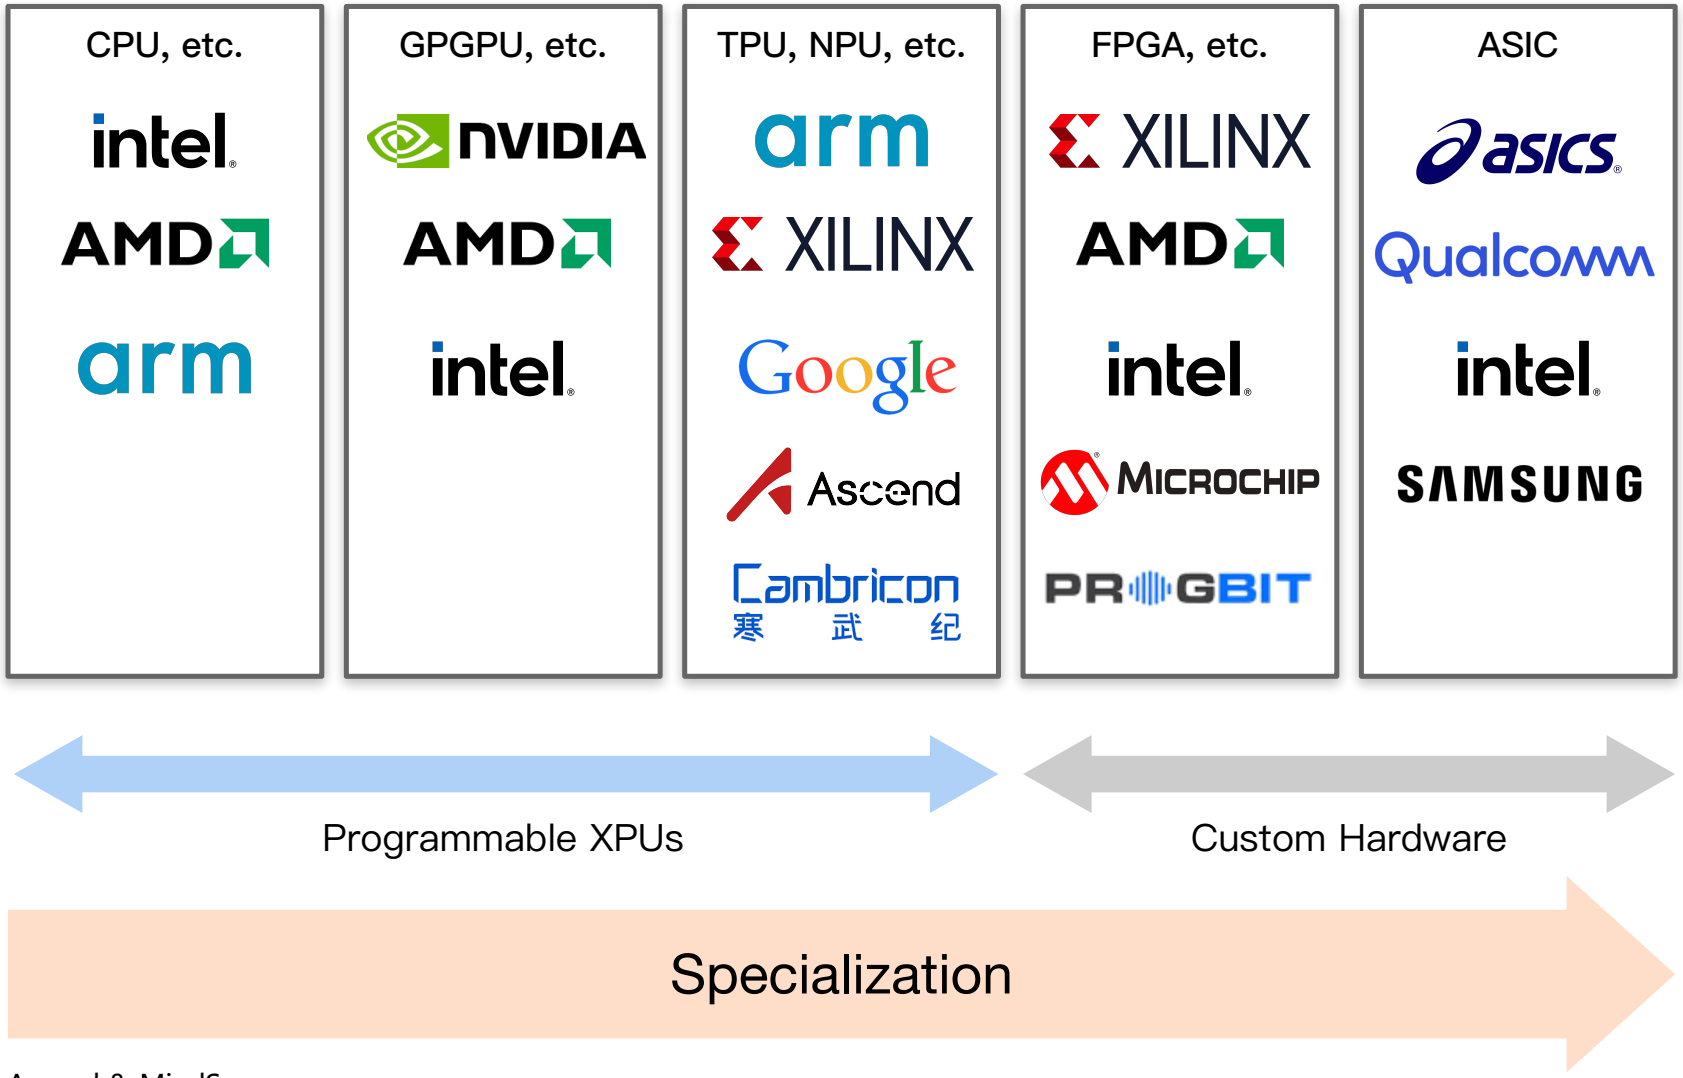

www.hiascend.com
www.mindspore.cn

How do we compile for this?

cādence

⇒ Not very compatible, inconsistent quality and scope

- … and don’t share much code

www.hiascend.com
www.mindspore.cn

How do accelerators work?

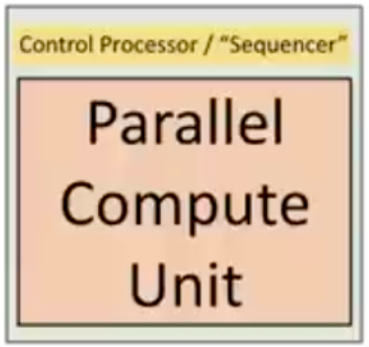

Control Processor / Sequencer

- Executes commands by the host driver app
- Handles booting and other housekeeping
- Diagnostics, security, debug, other functions

Some accelerators may do significantly more!

Ratio of control to parallel compute vary, as do the internal arch's of both

www.hiascend.com
www.mindspore.cn

Add a system interface

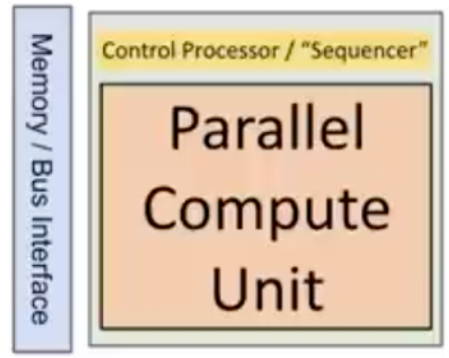

Communicate w other parts of the SoC, or to off-chip resources

Including DDR, HBM, … AMBA, PCI, CXL, etc depending on integration level

www.hiascend.com
www.mindspore.cn

“Oops we need some software”

Hardware

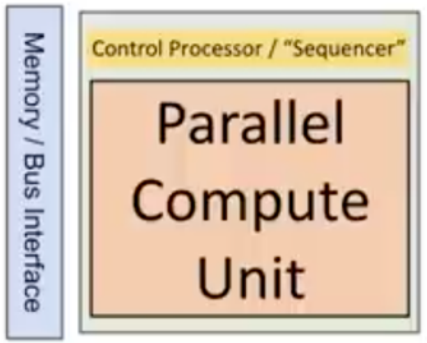

Software

Programming Model + Userspace API

Device Kernels

Control Proc Assembler + Kernel Driver

The SW people are called in after the accelerator is defined to “make it work”

www.hiascend.com
www.mindspore.cn

Larger accelerators go multicore/SIMT…

Hardware

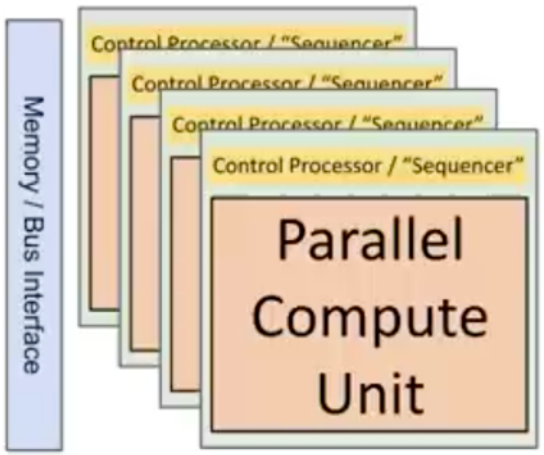

Software

Programming Model + Userspace API

Parallel Device Kernels

Control Proc Assembler + Kernel Driver

Use of more HW area is desired, requiring parallel control logic

www.hiascend.com
www.mindspore.cn

Tiling and heterogeneity for generality

Hardware

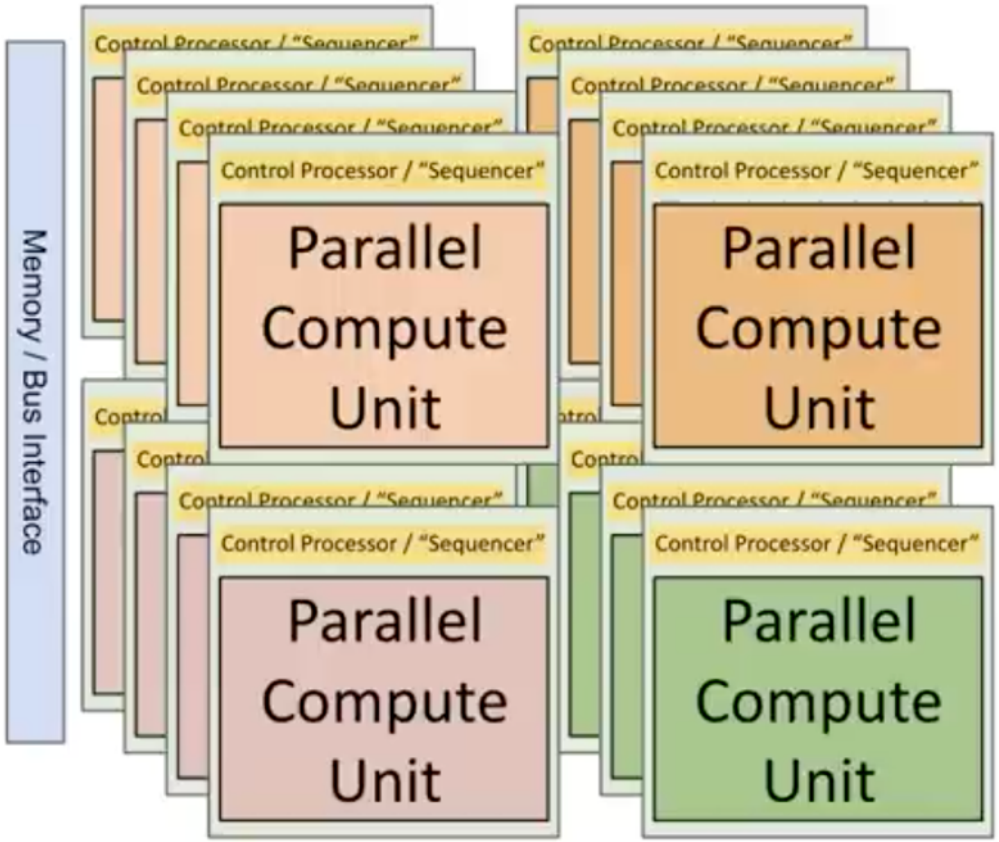

Software

Programming Model + Userspace API

Multistream Mgmt / Interop Parallelism
Memory + Communication Optimization
Heterogenous Device + Host fallback

Parallel Device Kernels

Control Proc Assembler + Kernel Driver

⇒ Also, hierarchical compute at the board, rack, and datacenter level

www.hiascend.com
www.mindspore.cn

Pro & Cons of hand written kernels

Benefits:

- Easy to get started, ability to get peak performance, hackability

Problem: **hand written kernels don’t scale**

- Expensive to maintain a library of 100’s to 1000’s of kernels
- Don’t scale to configurable IPs, not even memory hierarchy dimensions
- Don’t scale to device families, or evolving μarch’s over time
- Eventually end up limiting HW design space exploration / evolution

Often addressed with metaprogramming (aka “mini compilers”)

www.hiascend.com
www.mindspore.cn

“DSA Compilers” to the rescue

Hardware

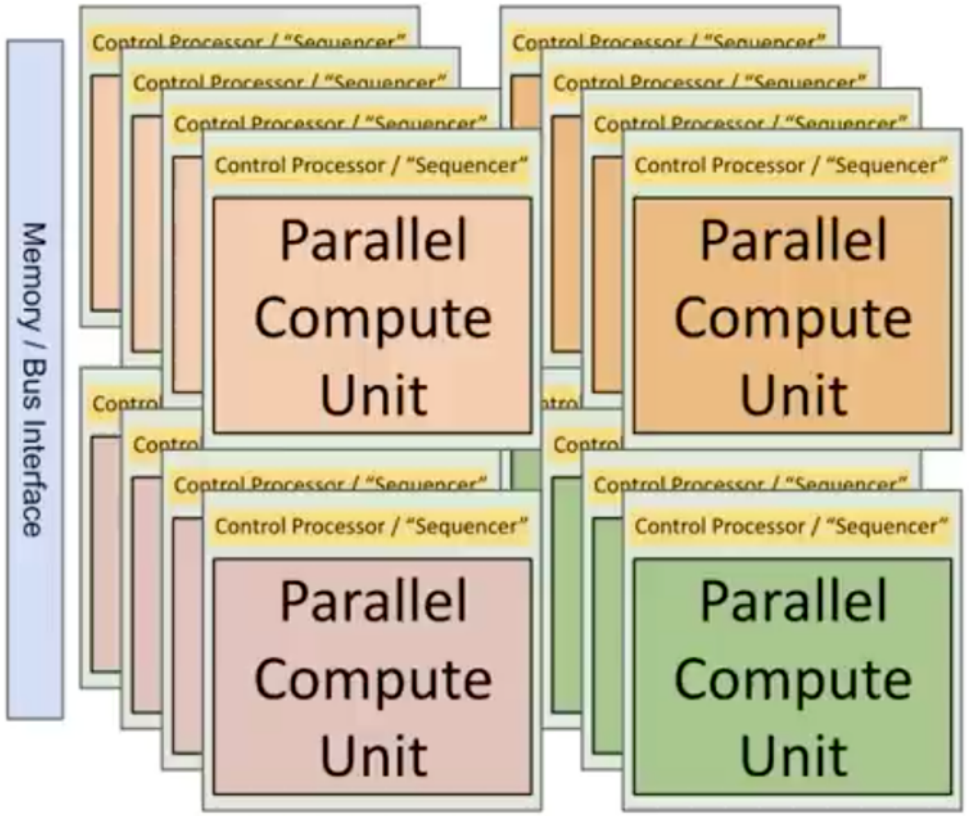

Software

Programming Model + Userspace API

Accelerator Kernel Compiler

Multistream Mgmt / Interop Parallelism
Memory + Communication Optimization
Heterogenous Device + Host fallback
Kernel Code Generation

Control Proc Assembler + Kernel Driver

www.hiascend.com
www.mindspore.cn

This is Hard!

… And we keep reinventing it over and over again.

… at the expense of usability and quality.

www.hiascend.com
www.mindspore.cn

Standardize your base Software

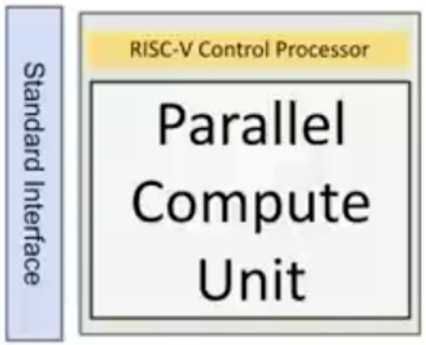

Write your kernels in C or LLVM IR!

- Use *existing* code generators
- Use *existing* simulators
- Step through them in a **debugger**

RISC-V **Compiler** + Kernel Drivers

www.hiascend.com
www.mindspore.cn

The next frontier: DSA Compilers?

"No one size fits all" compiler!

Shape of the problem is the same...
... but the accel details always vary

How do we get reuse?

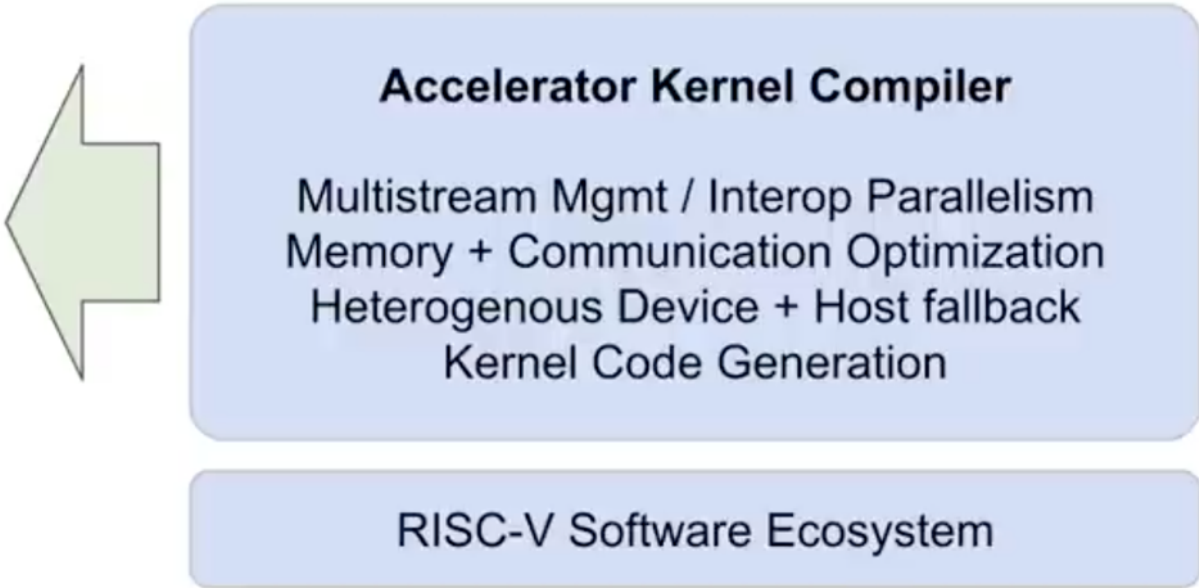

www.hiascend.com
www.mindspore.cn

RISC-V+MLIR: Uniting an Industry

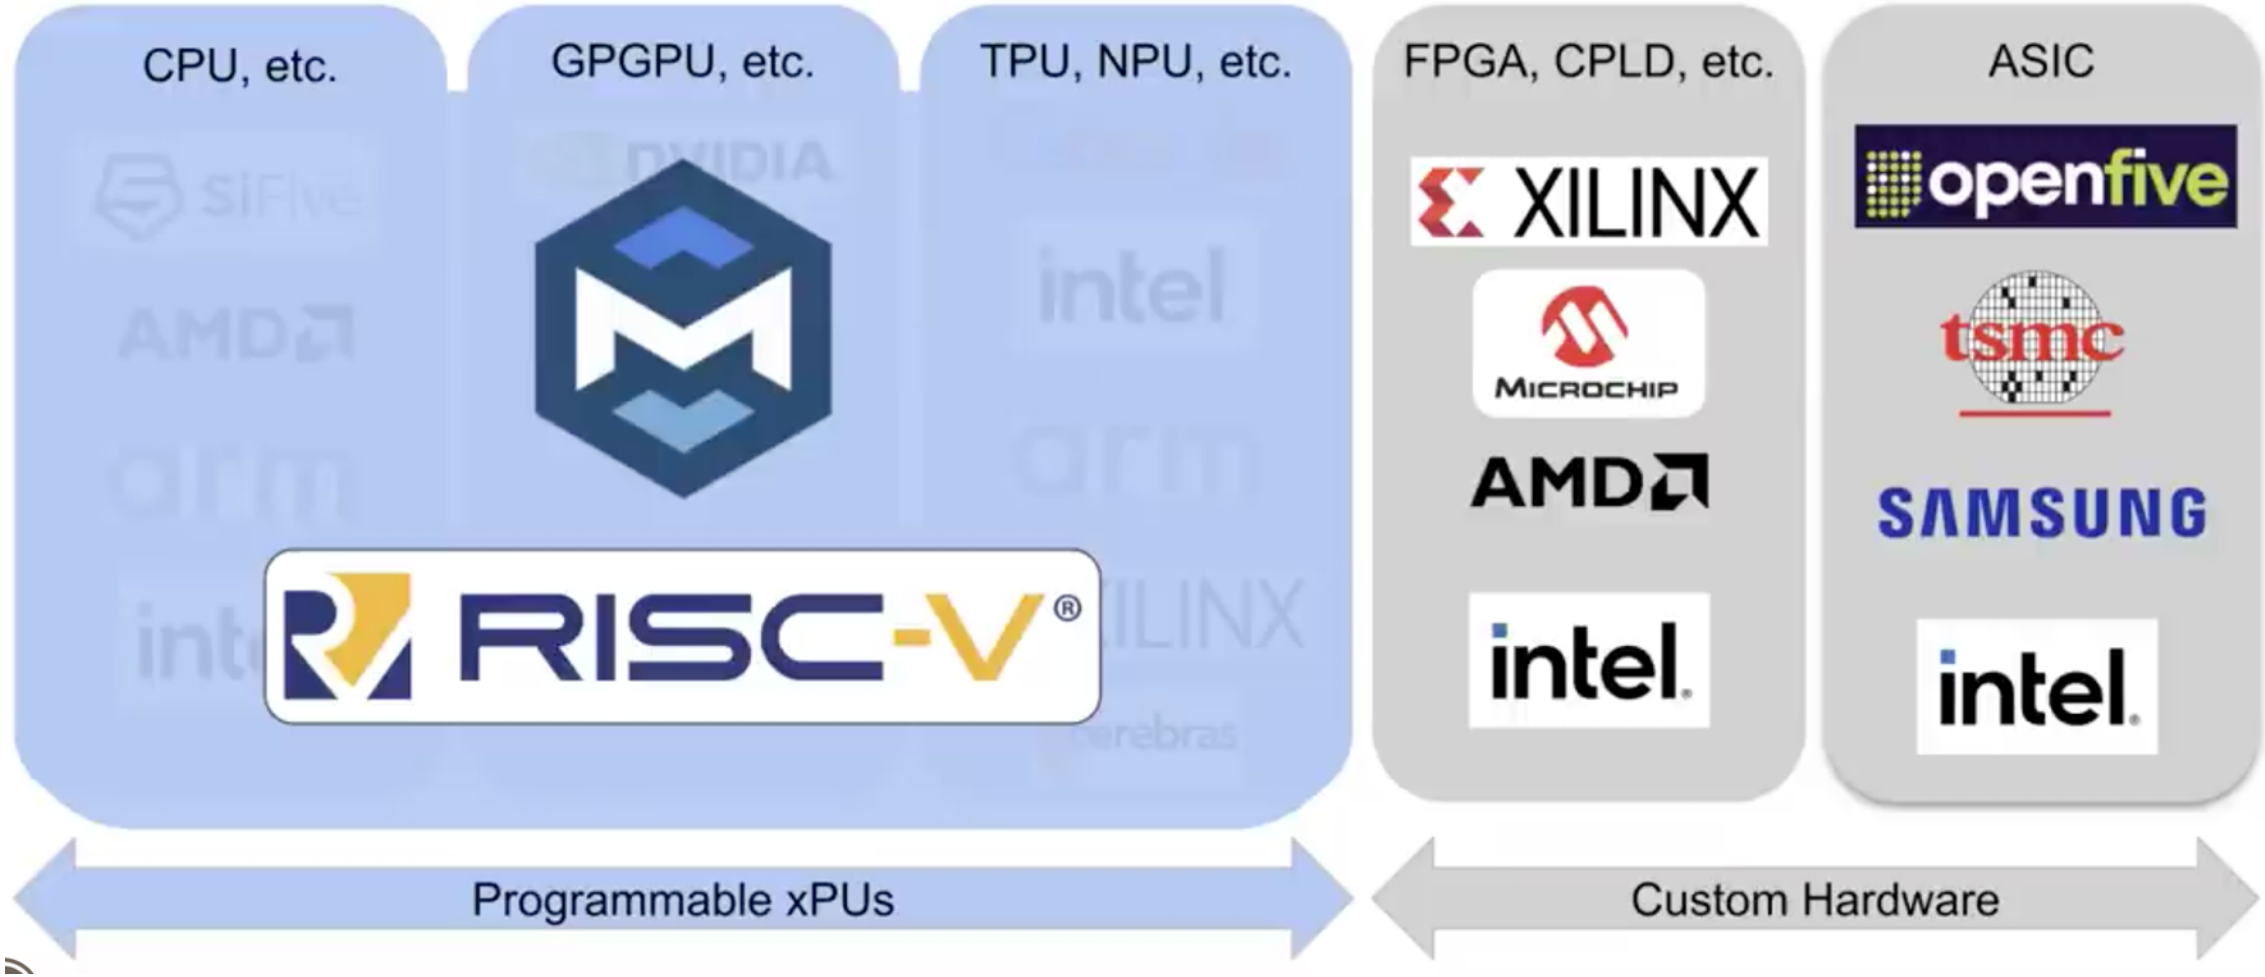

www.hiascend.com
www.mindspore.cn

Deep Learning is widely used

CV

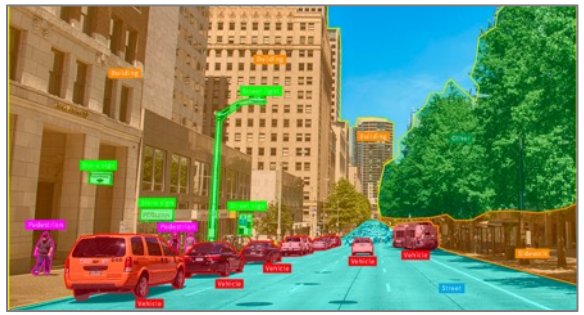

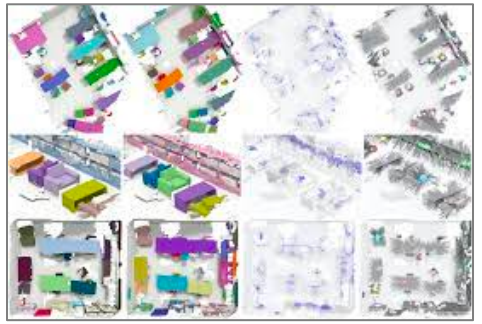

NLP

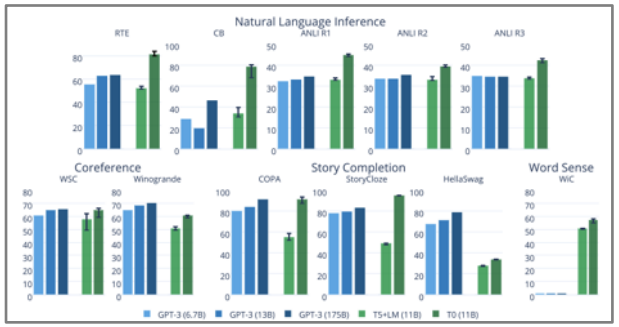

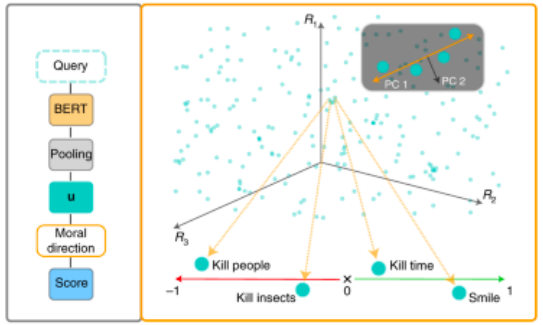

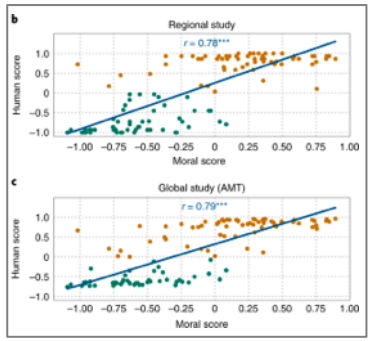

Science

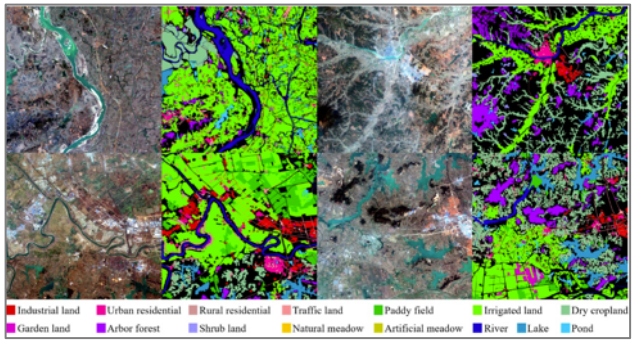

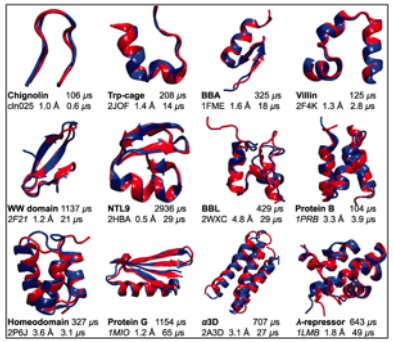

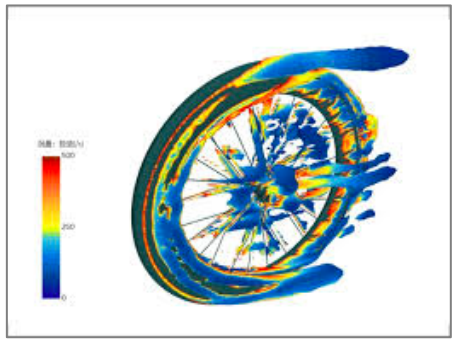

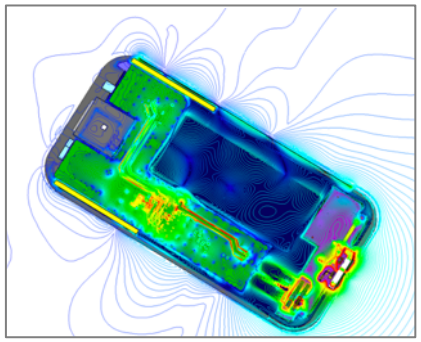

www.hiascend.com
www.mindspore.cn

AI Framework Problem

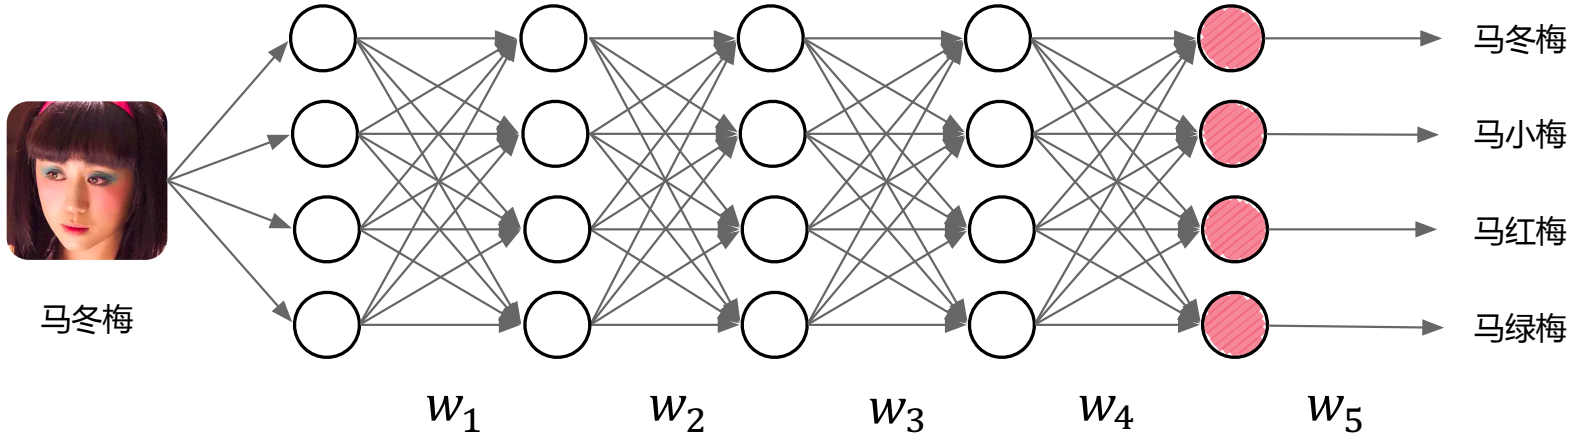

www.hiascend.com
www.mindspore.cn

AI Framework Problem

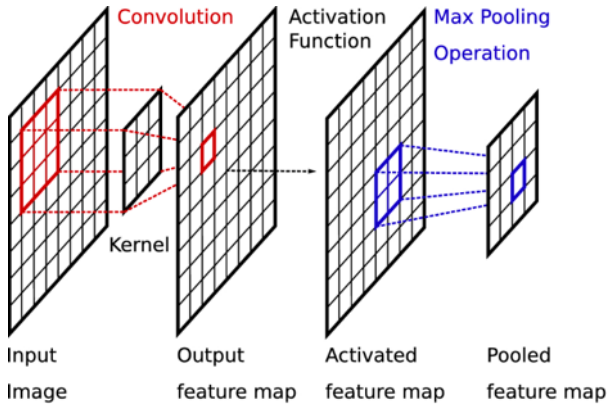

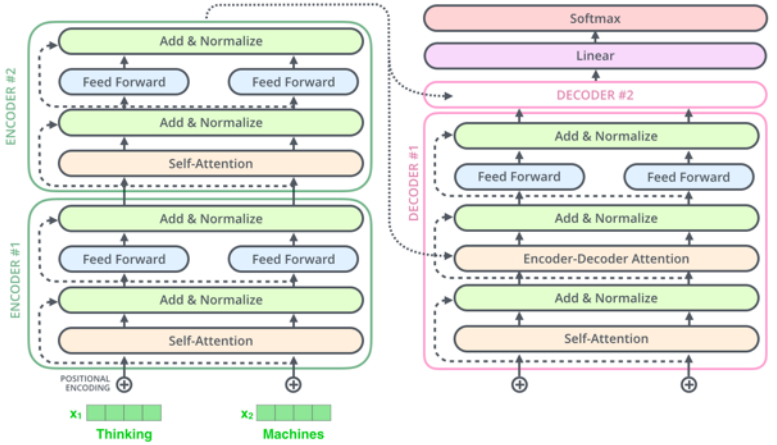

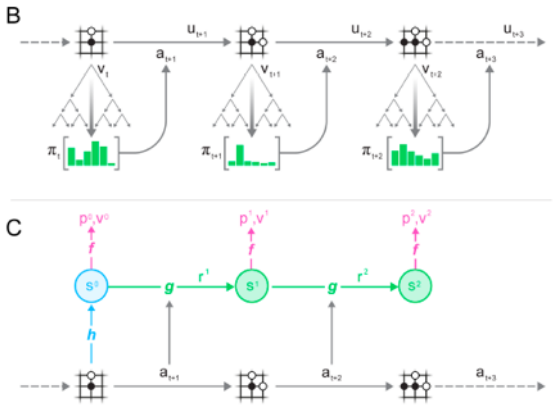

昇思 MindSpore

www.hiascend.com
www.mindspore.cn

The Gap Between SW and HW is huge!

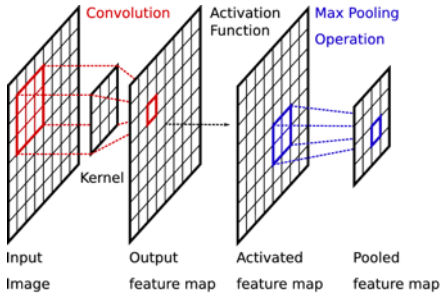

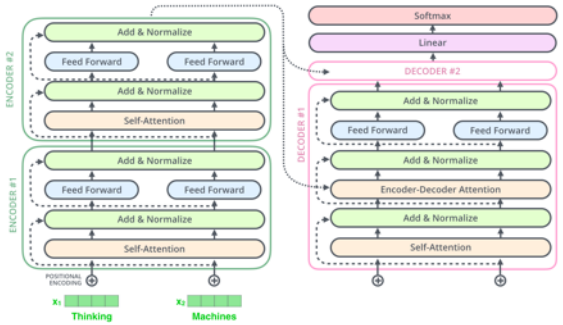

MindSpore TensorFlow

OneFlow 飞桨 PyTorch

www.hiascend.com
www.mindspore.cn

C Compilers leading into the early 90s

www.hiascend.com
www.mindspore.cn

Challenge

1. 越来越多新算子被提出，算子库的开发、维护、优化和测试工作量指数上升；
2. 专用加速芯片爆发导致性能可移植性成为一种刚需；

www.hiascend.com
www.mindspore.cn

Challenge: Operator

- 越来越多新算子被提出，算子库的开发、维护、优化和测试工作量指数上升：
  - 增加新算子，硬件不仅需要实现，还需要结合硬件进行特性优化和测试，尽量充分发挥硬件性能。以 Convolution 运算为例，需要将 Convolution 操作转换为 GEMM 矩阵乘法；算法新提出新的 Swish 算子，硬件需要新增 Swish 对应实现。
  - 硬件供应商还会有针对性发布优化库（如 MKL-DNN 和 CuDNN）。但是对于专用硬件，需要提供开发类似的优化库，不仅会增加大量算子优化、封装的工作，还会过于依赖库无法有效利用专用硬件芯片能力。

www.hiascend.com
www.mindspore.cn

Challenge: Optimization

- 专用加速芯片爆发导致性能可移植性成为一种刚需：
  - 大多数 NPU 使用 ASIC，在神经网络场景对计算、存储和数据搬运 做了特殊的指令优化，使得对 AI 相关的计算会提升性能，如 NVIDIA 虽然作为 GPGPU，但是 DGX 系列提供专用的 Tensor Code。
  - 不同厂商提供 XPU 的 ISA（Instruction Set Architecture）千奇百怪，一般缺乏如 GCC、LLVM 等编译工具链，使得针对 CPU 和 GPU已有的优化算子库和针对语言的优化 Pass 很难短期移植到 NPU 上。

www.hiascend.com
www.mindspore.cn

无限的算力和有限的精力

古有精卫填海

我想用钱试试

18岁的吴彦祖

你说我像苏大强？

www.hiascend.com
www.mindspore.cn

RISC-V+MLIR: Uniting an Industry

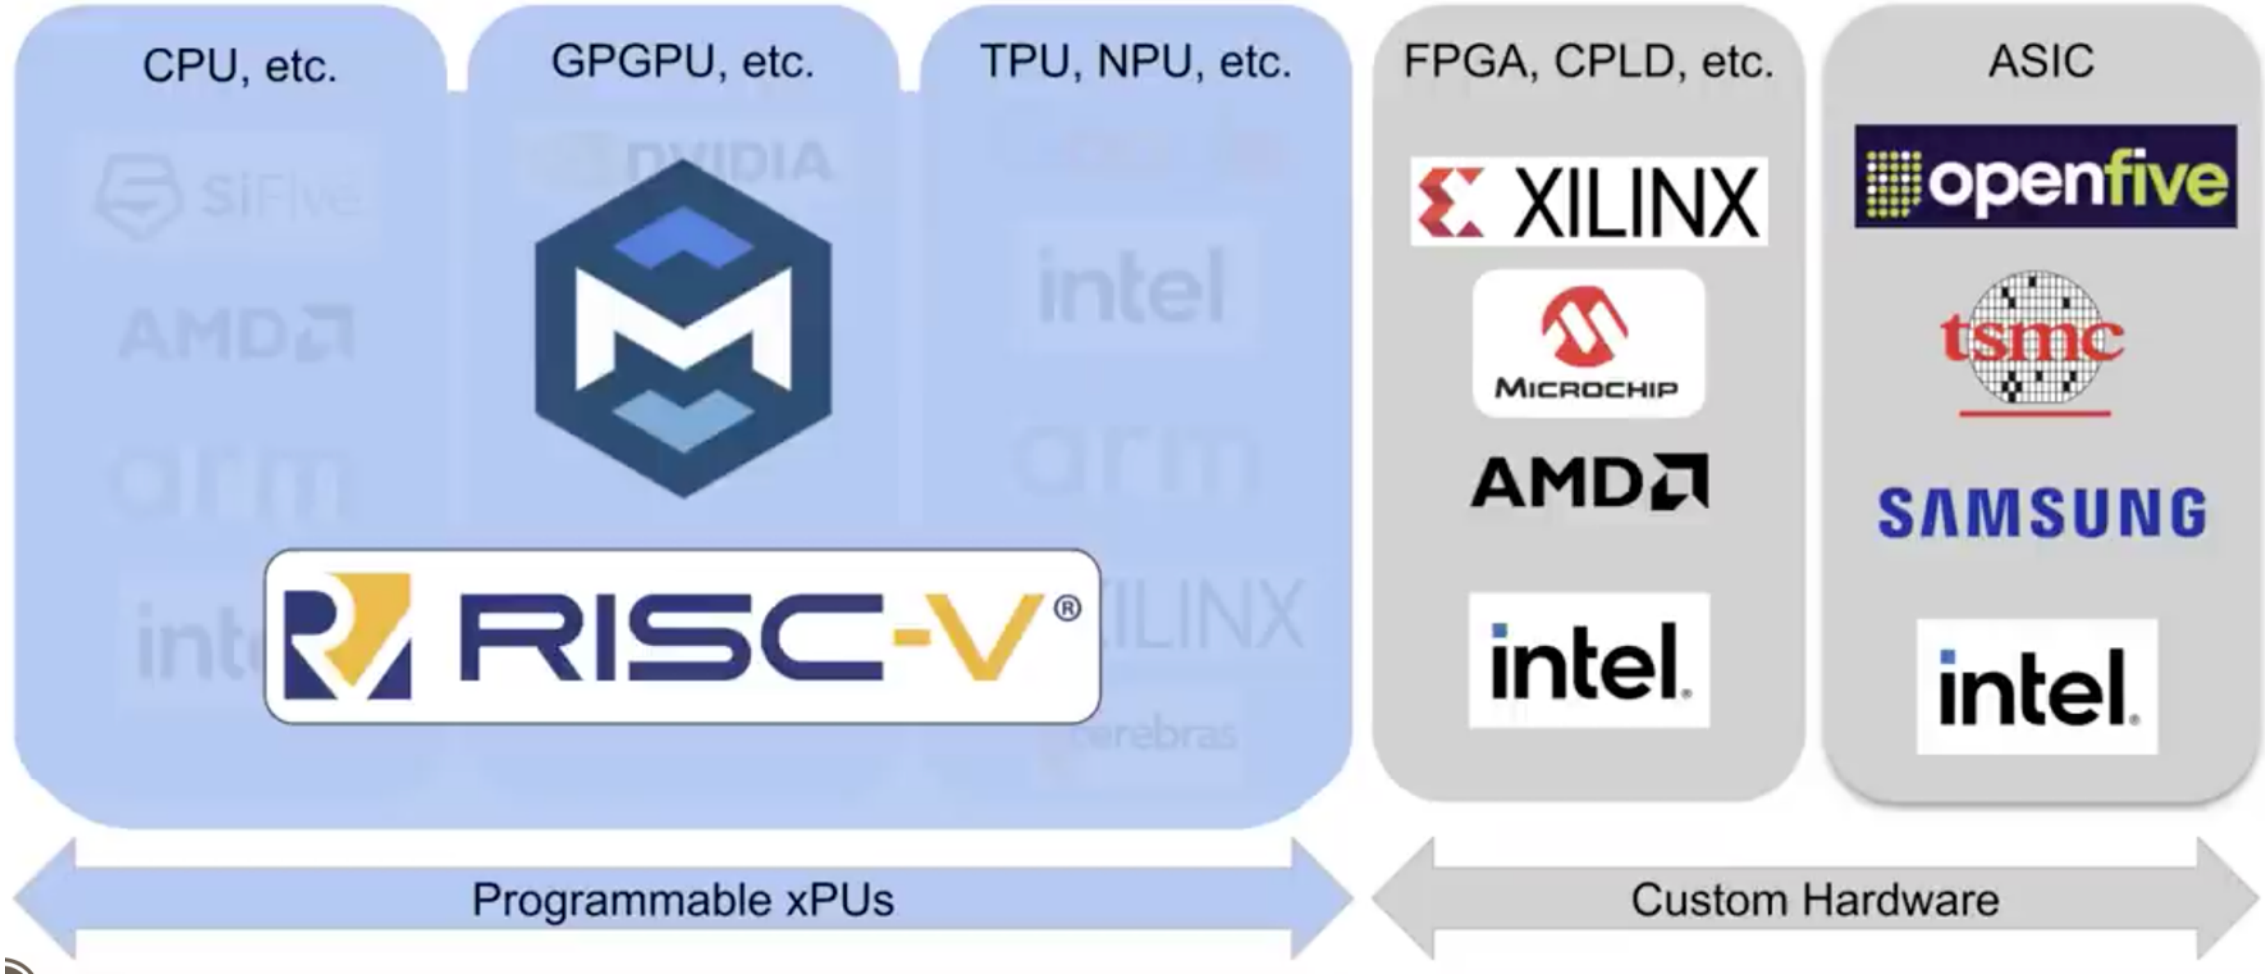

www.hiascend.com
www.mindspore.cn

传统编译器与AI编译器

AI 编译器跟传统编译器有什么区别吗？

www.hiascend.com
www.mindspore.cn

Likewise

- **目标相同：**通过自动化方式进行程序优化和代码生成，从而降低对不同硬件的手工优化；
- **优化方式类似：**在编译优化层通过统一 IR 执行不同的Pass进行优化，从而提高执行性能；
- **软件结构栈类似：**分成前端、优化、后端三段式，IR 解耦前端和后端使得模块化表示；
- **AI编译器依赖传统编译器：**AI编译器对 Graph IR 进行优化后，将优化后的 IR 转换成传统编译器 IR，最后依赖传统编译器针进行机器码生成。

www.hiascend.com
www.mindspore.cn

Difference

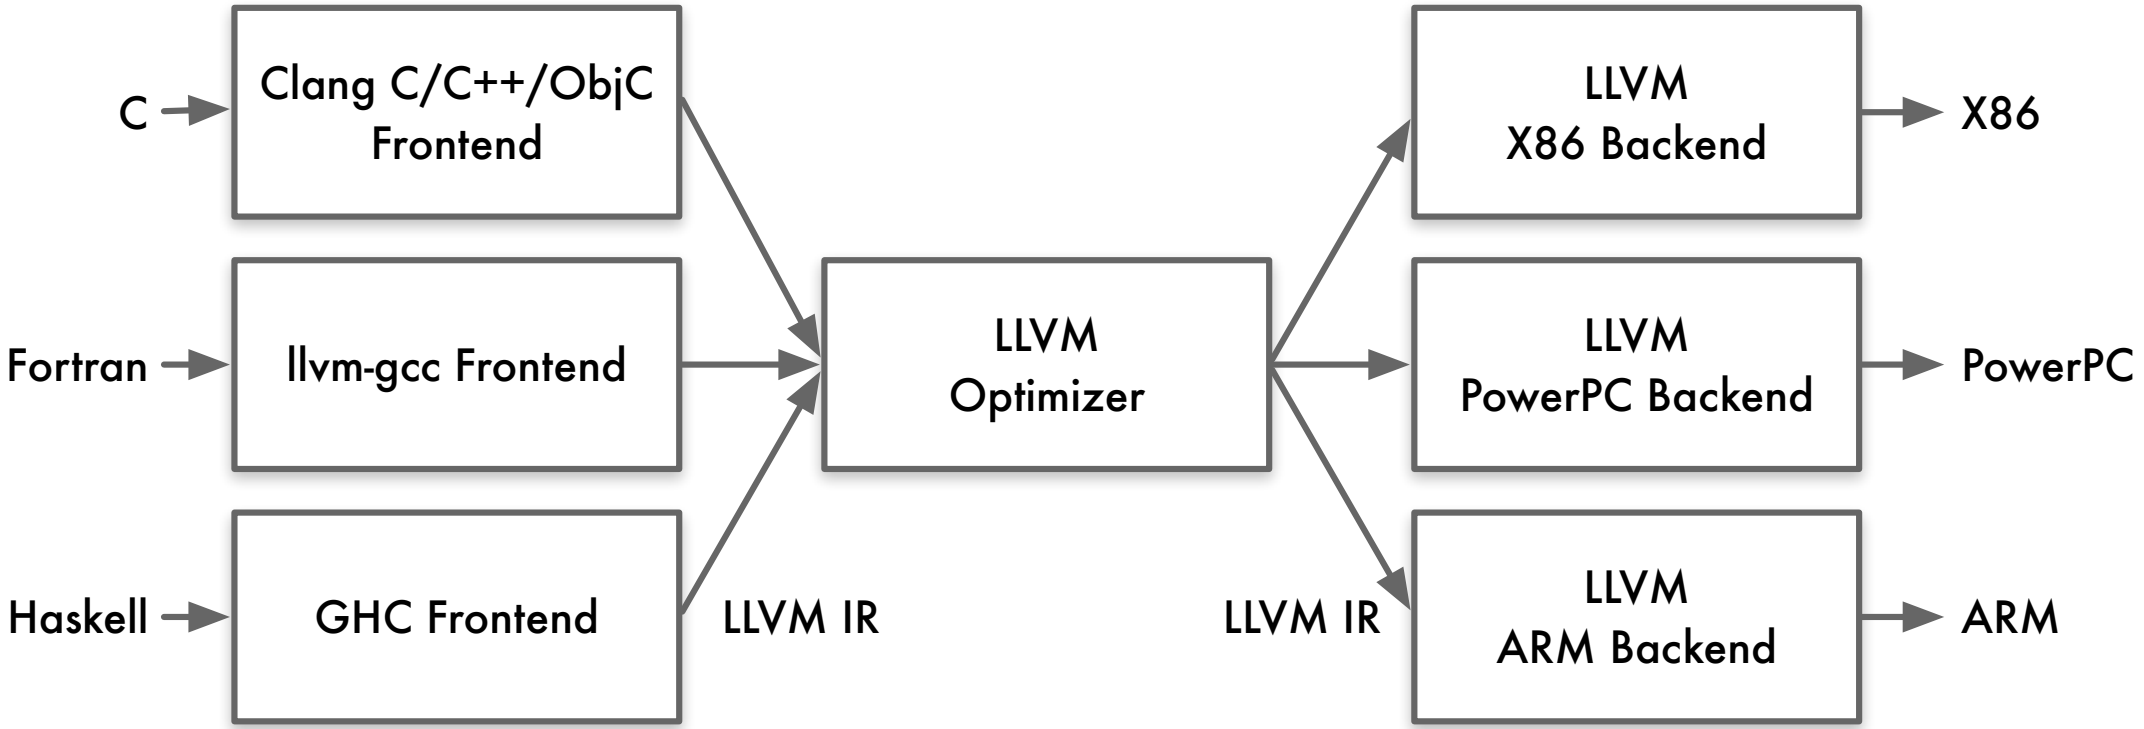

www.hiascend.com
www.mindspore.cn

Difference

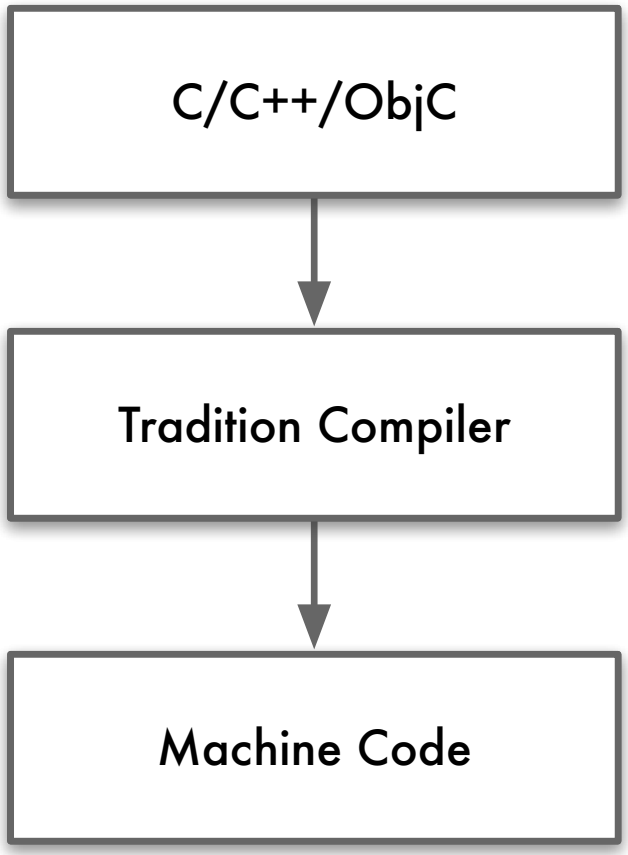

输入高级语言输出低级语言

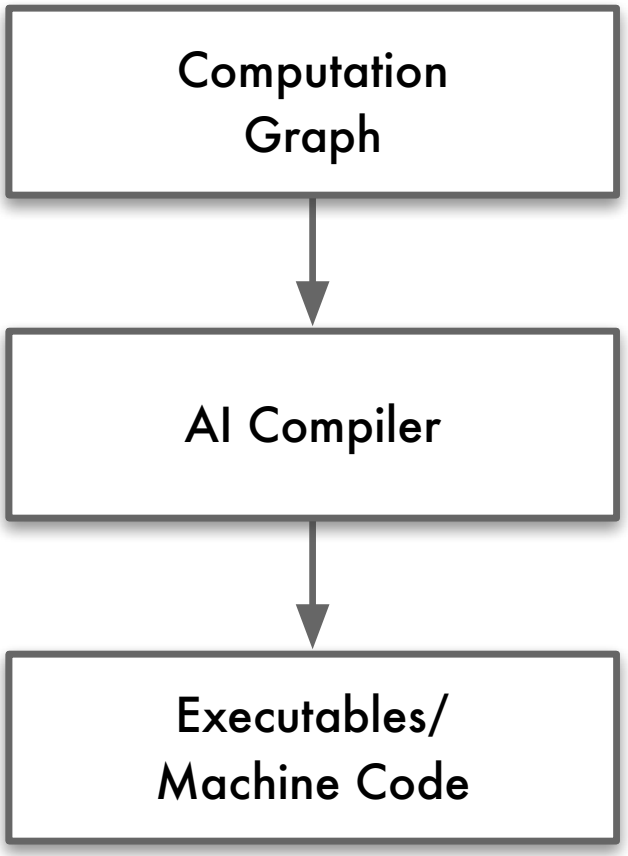

输入计算图/算子，输出低级语言

www.hiascend.com
www.mindspore.cn

Difference

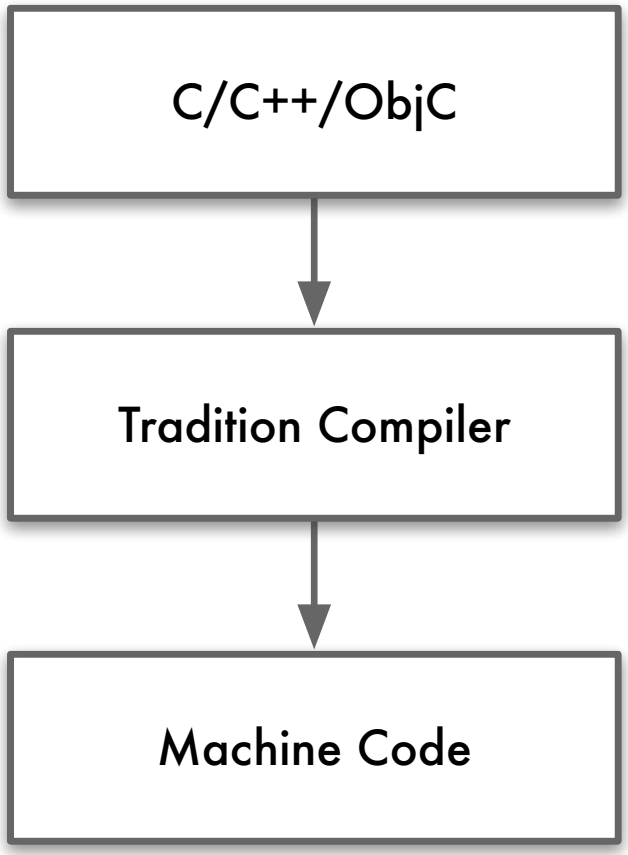

主要问题是降低编程难度

其次是优化程序性能

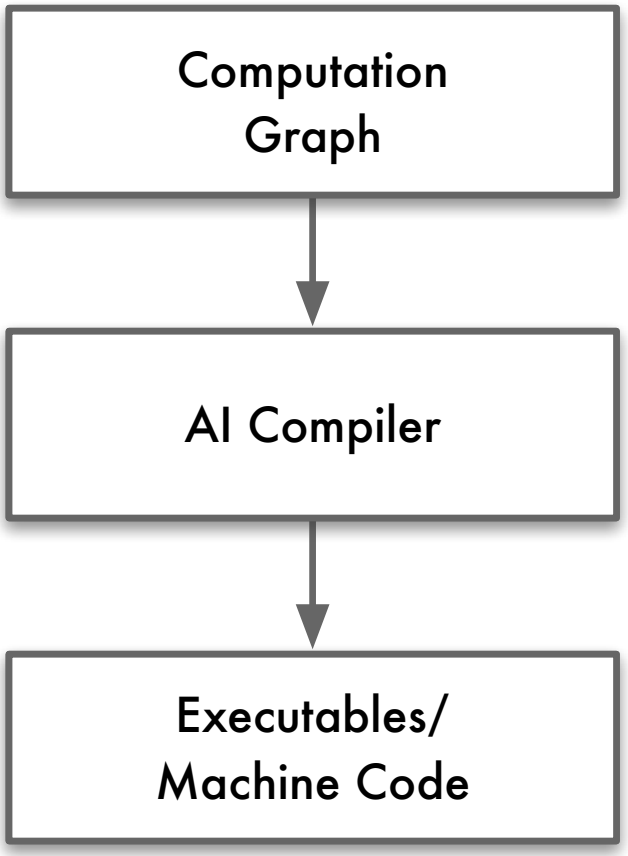

主要优化程序性能

其次是降低编程难度

www.hiascend.com
www.mindspore.cn

Difference

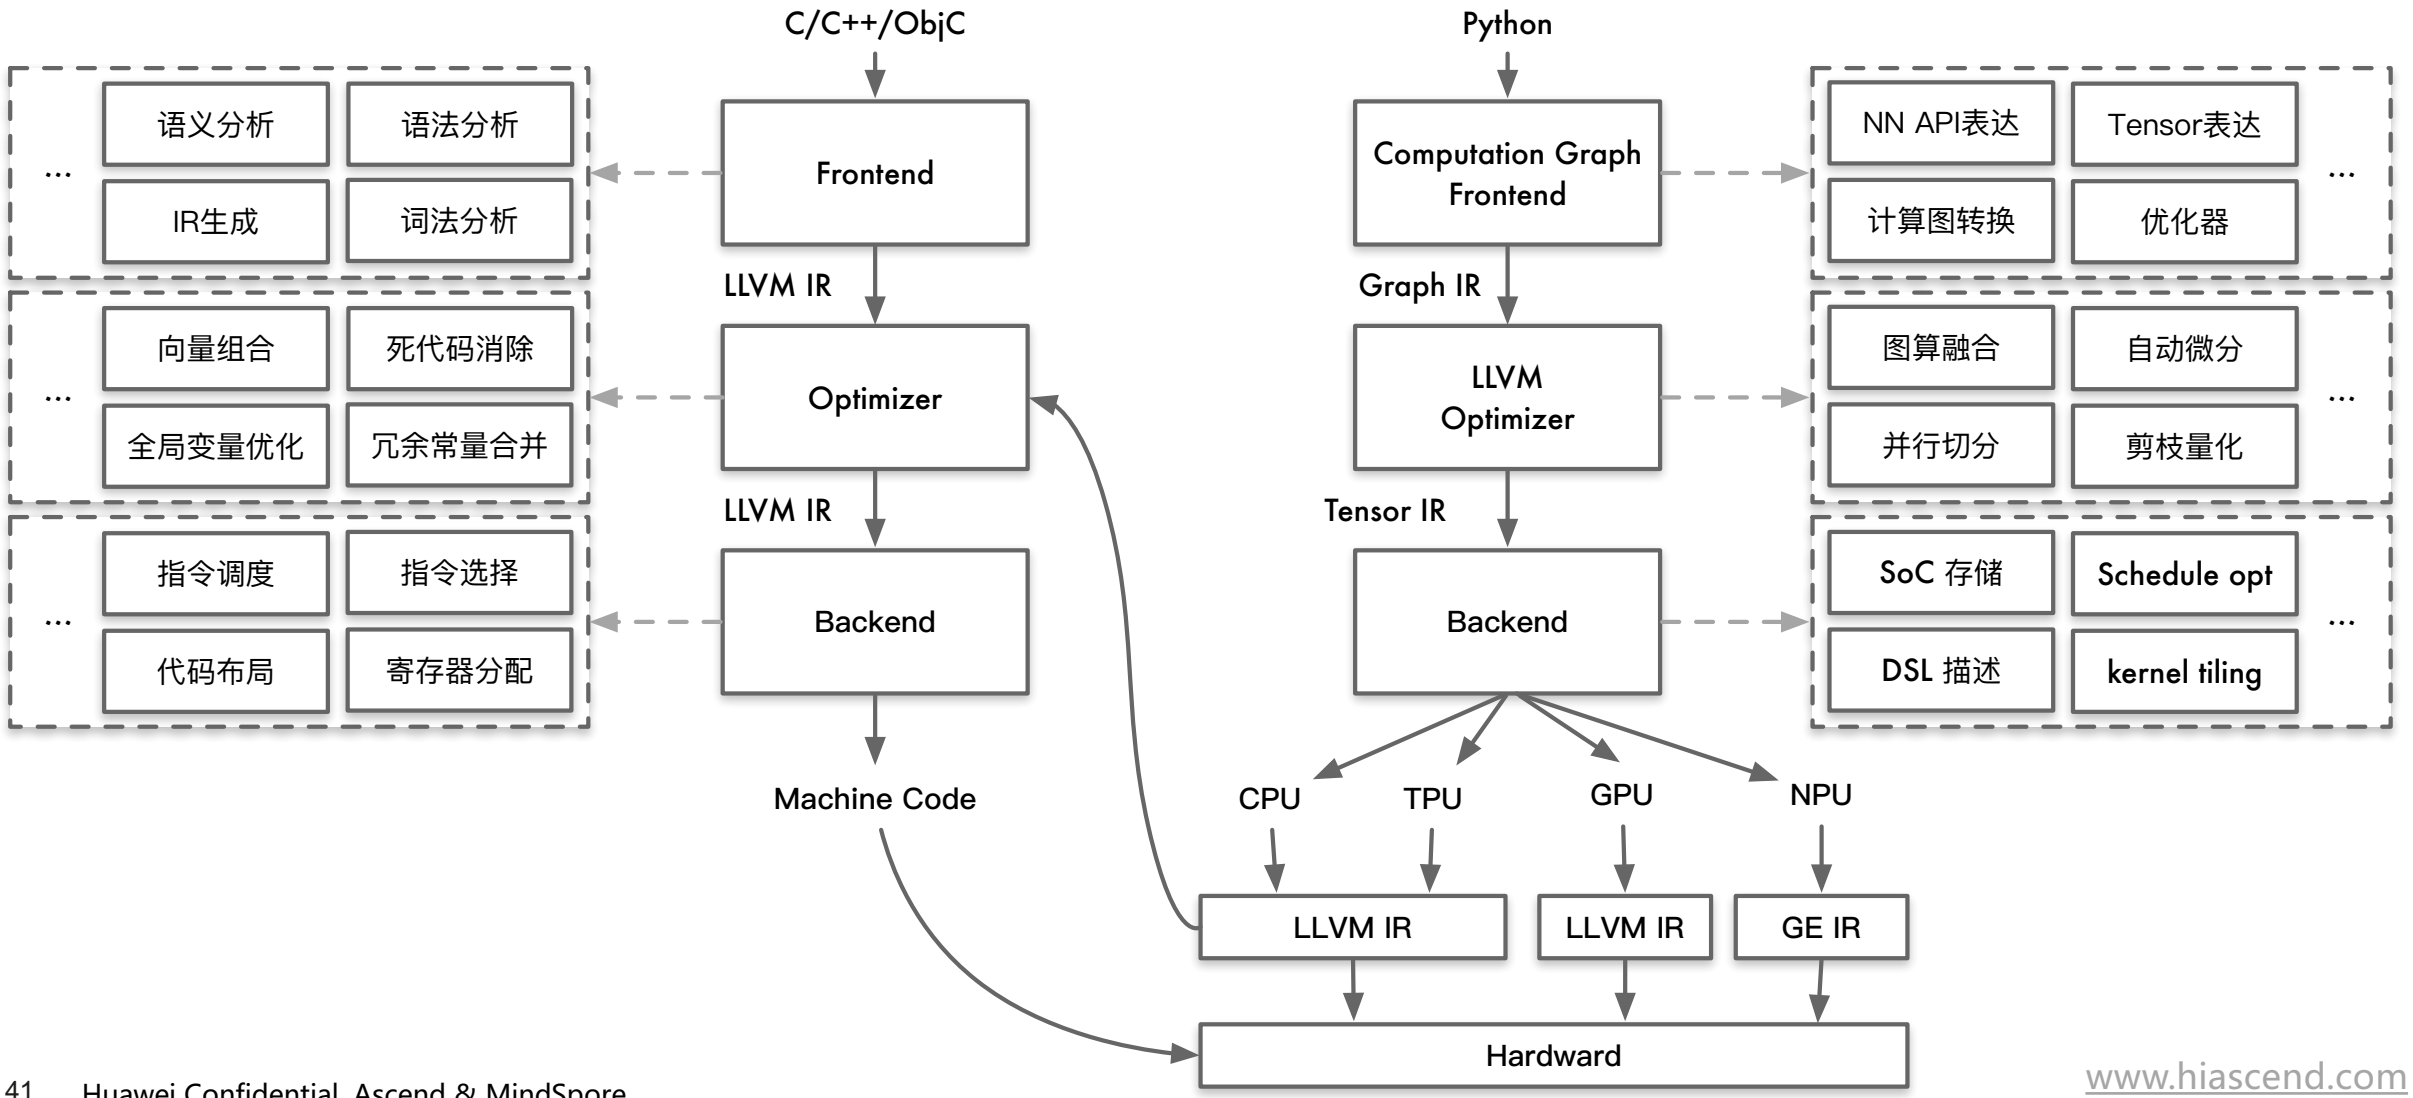

www.hiascend.com
www.mindspore.cn

Difference

- **IR 差异：**AI 编译器的 IR 与传统编译器的IR所抽象出来的概念和意义并不相同。
  - AI编译器一般会有 high-level IR，用来抽象描述深度学习模型中的运算，如：Convolution、Matmul 等，甚至部分会有 Transformer 带有图的结构。
  - 传统编译器相对而言 low-level IR，用于描述基本指令运算，如 load、store 等。有了high-level IR，AI编译器在描述深度学习模型类 DSL 更加方便。

www.hiascend.com
www.mindspore.cn

Difference

- **优化策略：**AI 编译器面向AI领域，优化时引入更多领域特定知识，从而进行更 high-level，更加 aggressive 优化手段。如：
  - AI编译器在 high-level IR 执行算子融合，传统编译器执行类似 loop fusion 时候，往往更加保守。缺点是可能会导致调试执行信息跟踪难；
  - AI编译器可以降低计算精度，比如int8、fp16、bf16等，因为深度学习对计算精度不那么敏感。但传统编译器一般不执行改变变量类型和精度等优化。

www.hiascend.com
www.mindspore.cn

Copyright©2014 Huawei Technologies Co., Ltd. All Rights Reserved.
The information in this document may contain predictive statements including, without limitation, statements regarding the future financial and operating results, future product portfolio, new technology, etc. There are a number of factors that could cause actual results and developments to differ materially from those expressed or implied in the predictive statements. Therefore, such information is provided for reference purpose only and constitutes neither an offer nor an acceptance. Huawei may change the information at any time without notice.